## **K Nearest Neighbors**







#### **Non-Parametric Models: Intuition**

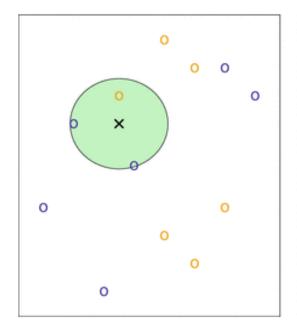
- A "parametric" model is one in which we make assumptions about the functional form of the model (e.g., a model is linear, errors are normally distributed, etc.).
- These assumptions **help** us define parameters (e.g., means, variance, coefficients) to **simplify** the model and its solutions.
- But parameters also constrain the solution to those that conform to these parameters (e.g., data points never touch the regression line)
- In contrast, a "non-parametric" model does not make such assumptions and the models simply try to get as close as possible to the data points.
- For example, if you find 10 teenagers in a store in the mall, you may be predicting that the next customer to walk in will be a teenager, without resorting to regression models or normal distributions.

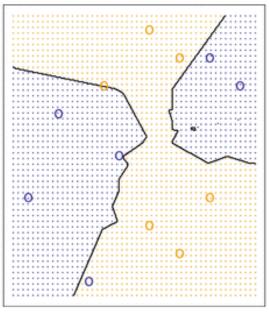




### K Nearest Neighbors (KNN): Intuition

- KNN is a non-parametric method most suitable for data in which we have very little knowledge about its distribution
- KNN simply classifies a test observation based on the classification of the K training closest training points.
- The resulting classification boundary for a K=3 KNN is illustrated below





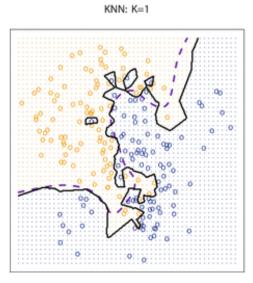


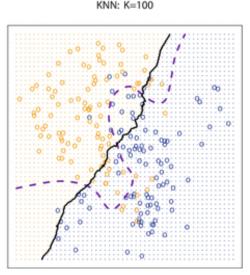


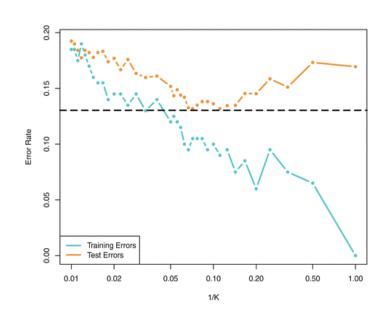
### Selecting K



- The solid line in the diagrams below illustrate KNN classification boundaries for K = 1 and K = 100
- As the left diagrams show, a larger K produce smoother classification boundaries; a smaller K produce better training accuracy, which naturally does not ensure test accuracy.
- As the diagram on the right illustrates, training error tends to decline when K increases (1/K decreases); but the test error bottoms out around K = 10 (1/K = 0.10)











knn () {class} → "K-Nearest Neighbors" function in the {class} package to fit KNN models. The {class} package has functions to fit various classification models.

One nice feature of the knn () function is that it does estimation and prediction in one step. You first need to define the training and test data sets:

train.x=cbind(x1,x2,etc.) [train,]  $\rightarrow$  Use the column bind function to create a predictor variable matrix; [train,] is the record index defined to select the training data set (see Machine Learning lecture)

test.x=cbind(x1,x2,etc.) [!train,]  $\rightarrow$  [!train,] Is the record index vector for all records not (!) in the train set, which we use for the test set.

knn.pred=knn (train.X, test.X, y, k=1)  $\rightarrow$  y is the classification variable; k=1 specifies 1 neighbor in the model; use k=n to specify n neighbors.





# KOGOD SCHOOL of BUSINESS

